

# Modern Operating System Concepts

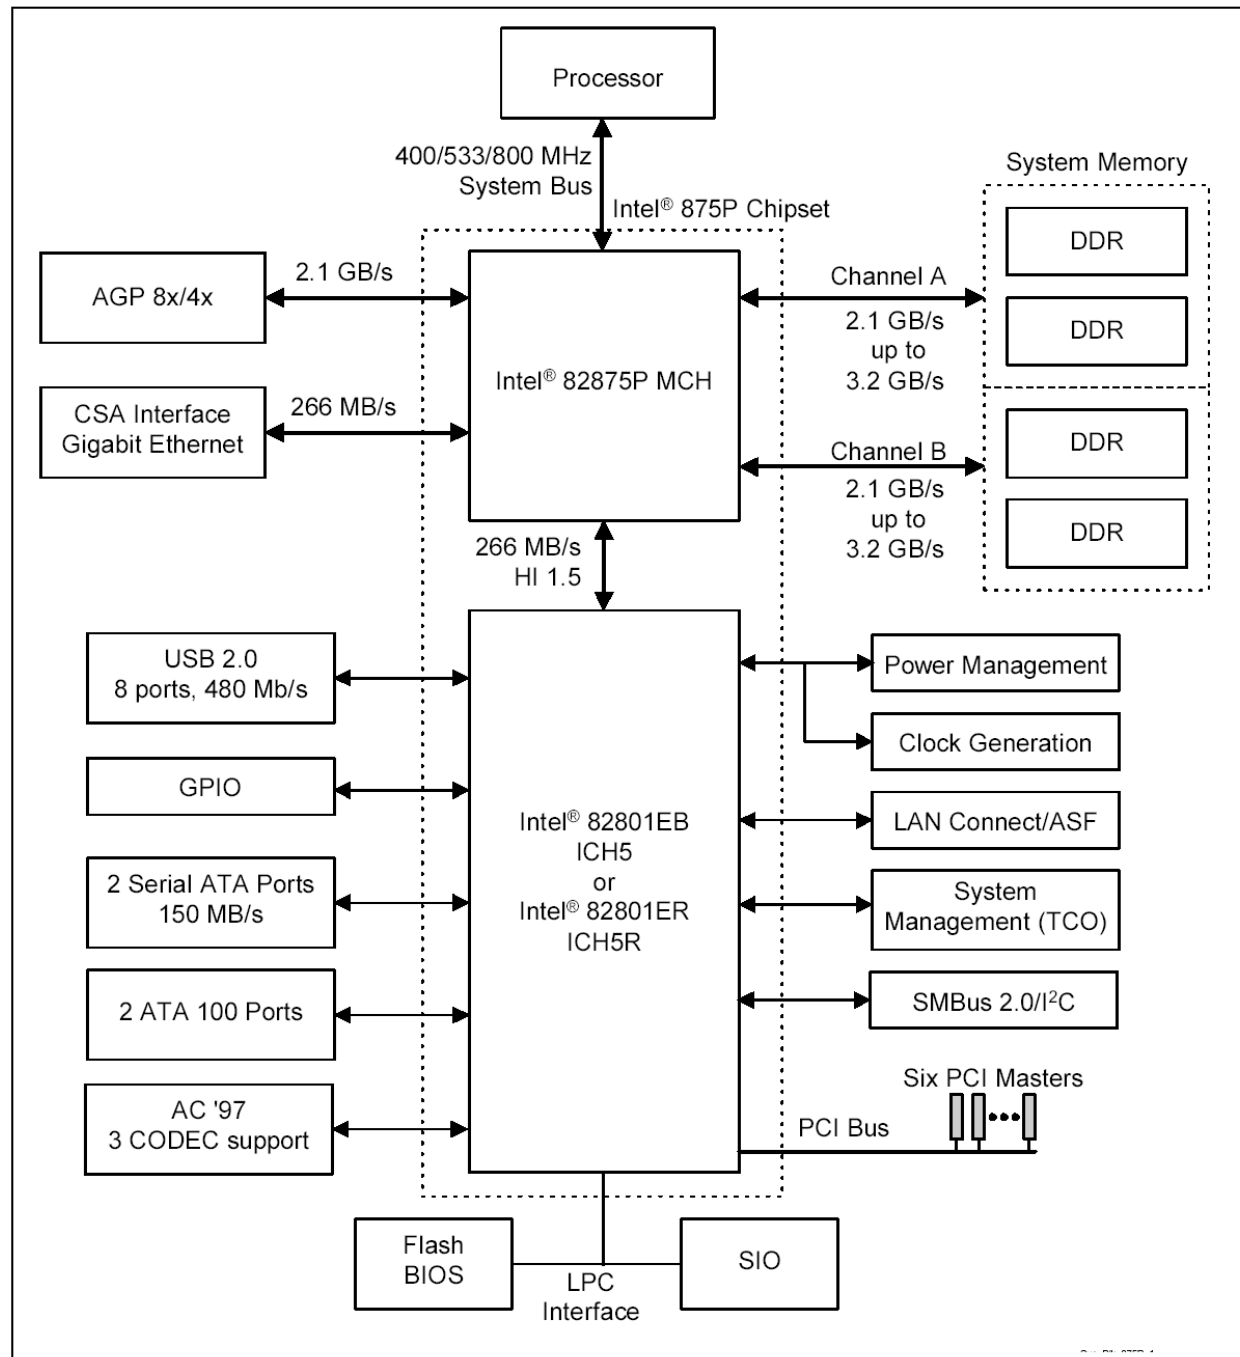
**Benjamin Brewster**

Adapted from slides by Jon Herlocker, OSU

# The Underlying Hardware

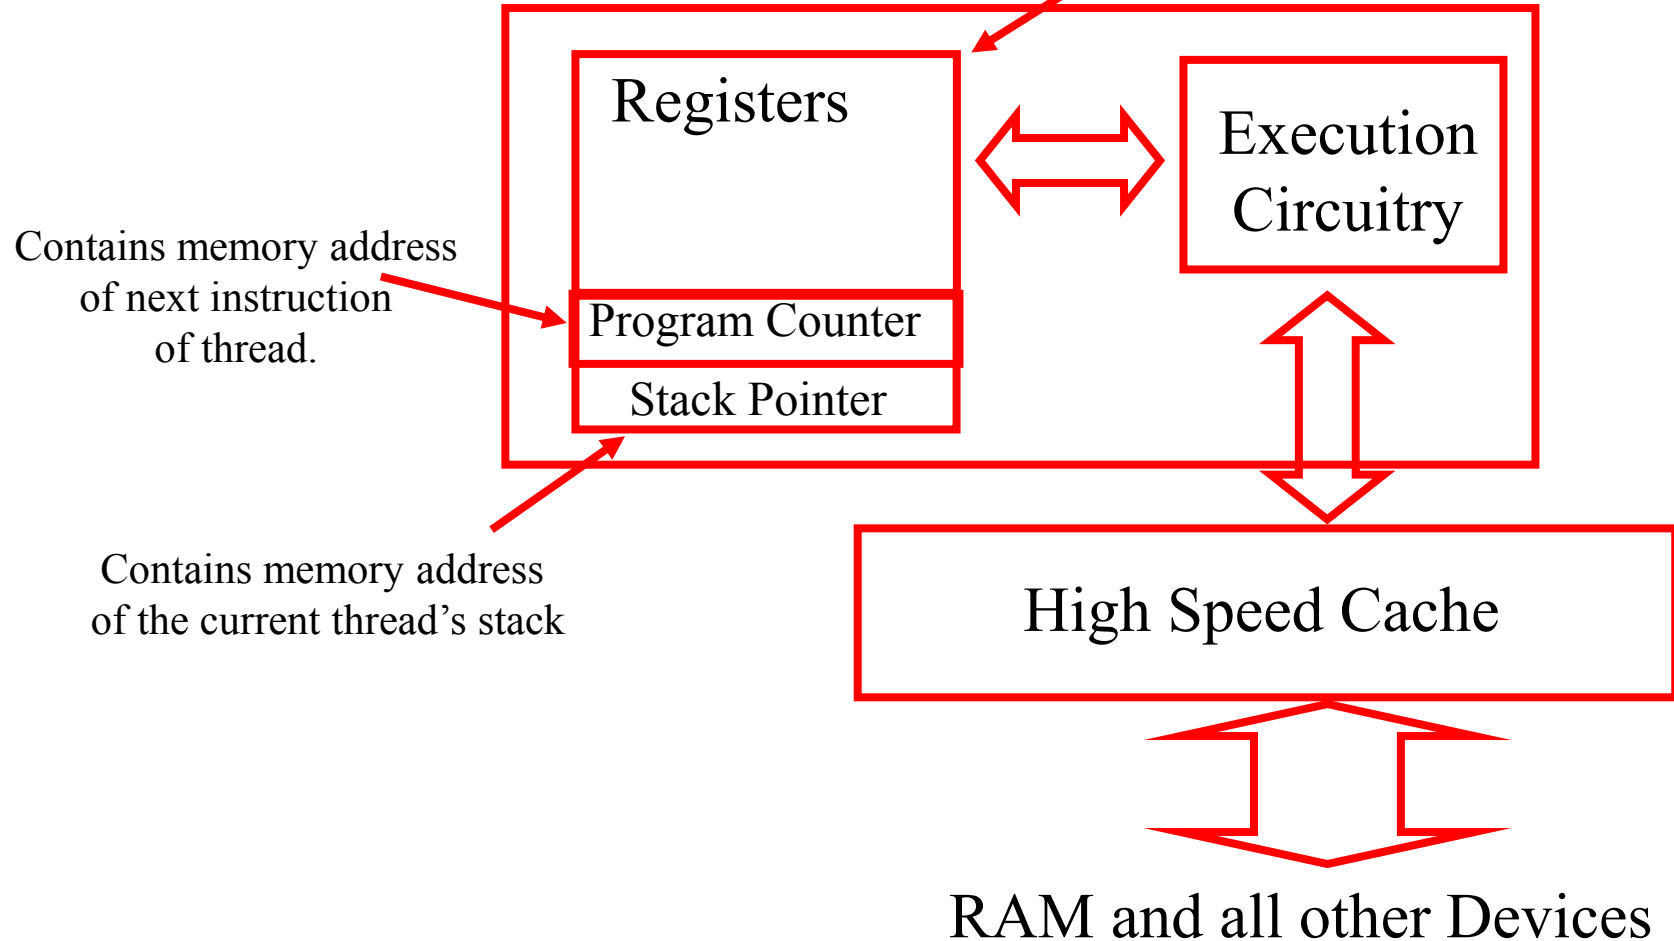
- The OS provides applications access to the hardware in an *abstracted* manner
- What does that hardware actually look like?

Figure 1. Intel® 875P Chipset System Block Diagram




# A Processor

Each register  
stores a number



# Memory

Memory is an array of *bytes*



Memory is temporary storage,  
Much slower than the processor  
and cache but much faster than a disk

# Other storage devices


- Persistent storage
  - Magnetic disk
  - Non-volatile memory (flash, etc.)
  - Magnetic tapes
  - Optical (CD-ROM, DVDs, BD, HD-DVD, etc.)
  - Slow, but persist without power

# Virtual Memory

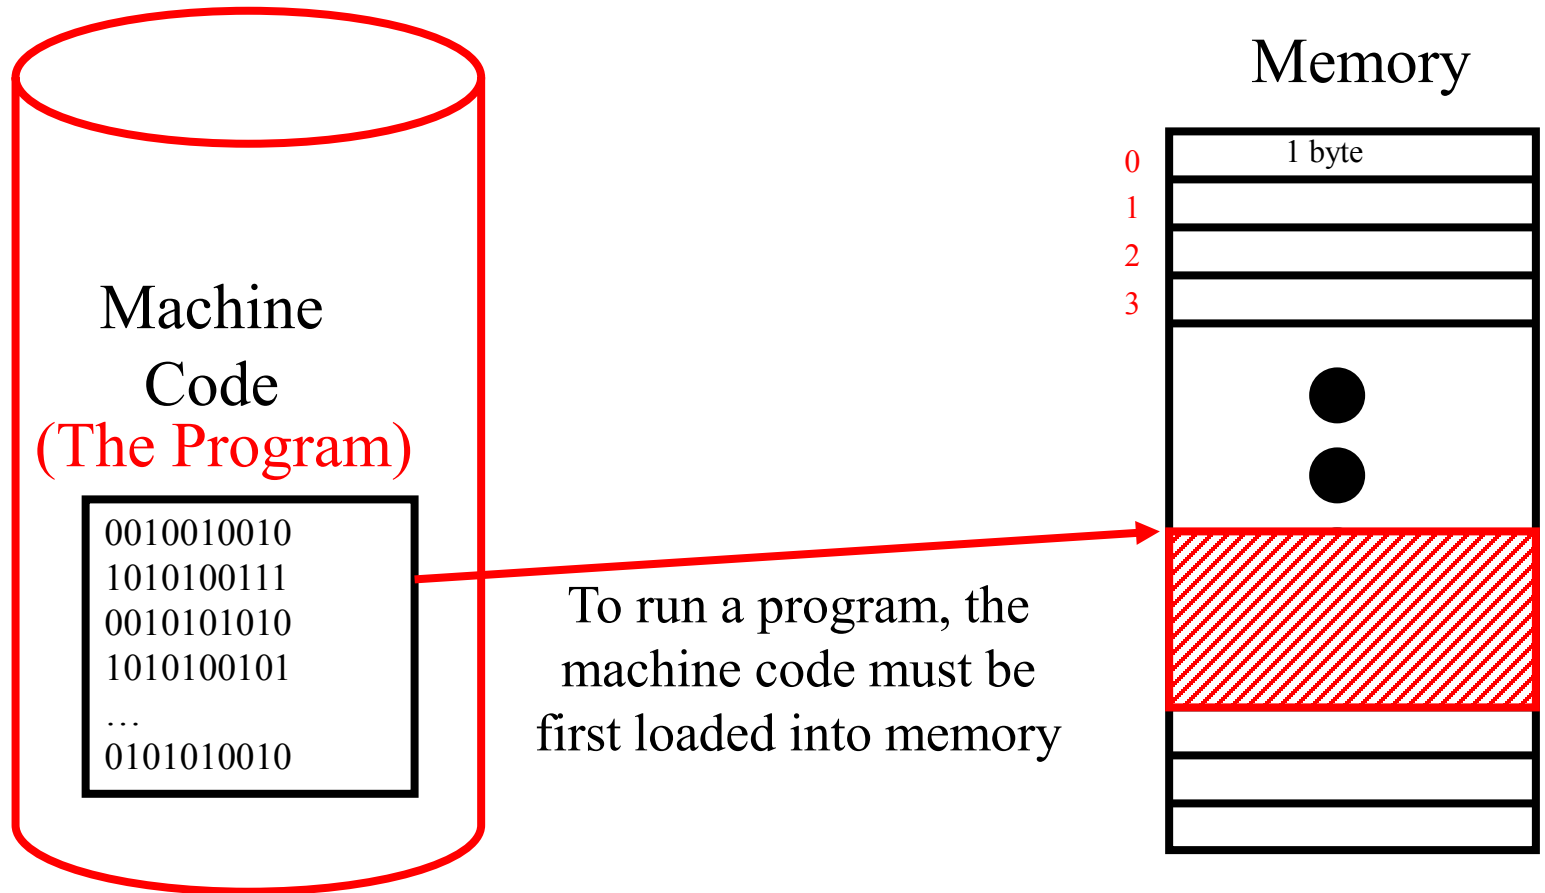
- Virtual memory hardware creates the illusion of
  - Exclusive memory
    - Not shared
  - Unlimited memory
    - (up to the maximum address size)

Thus a program's *address space* begins at 0

And ends at the largest possible address processor can handle

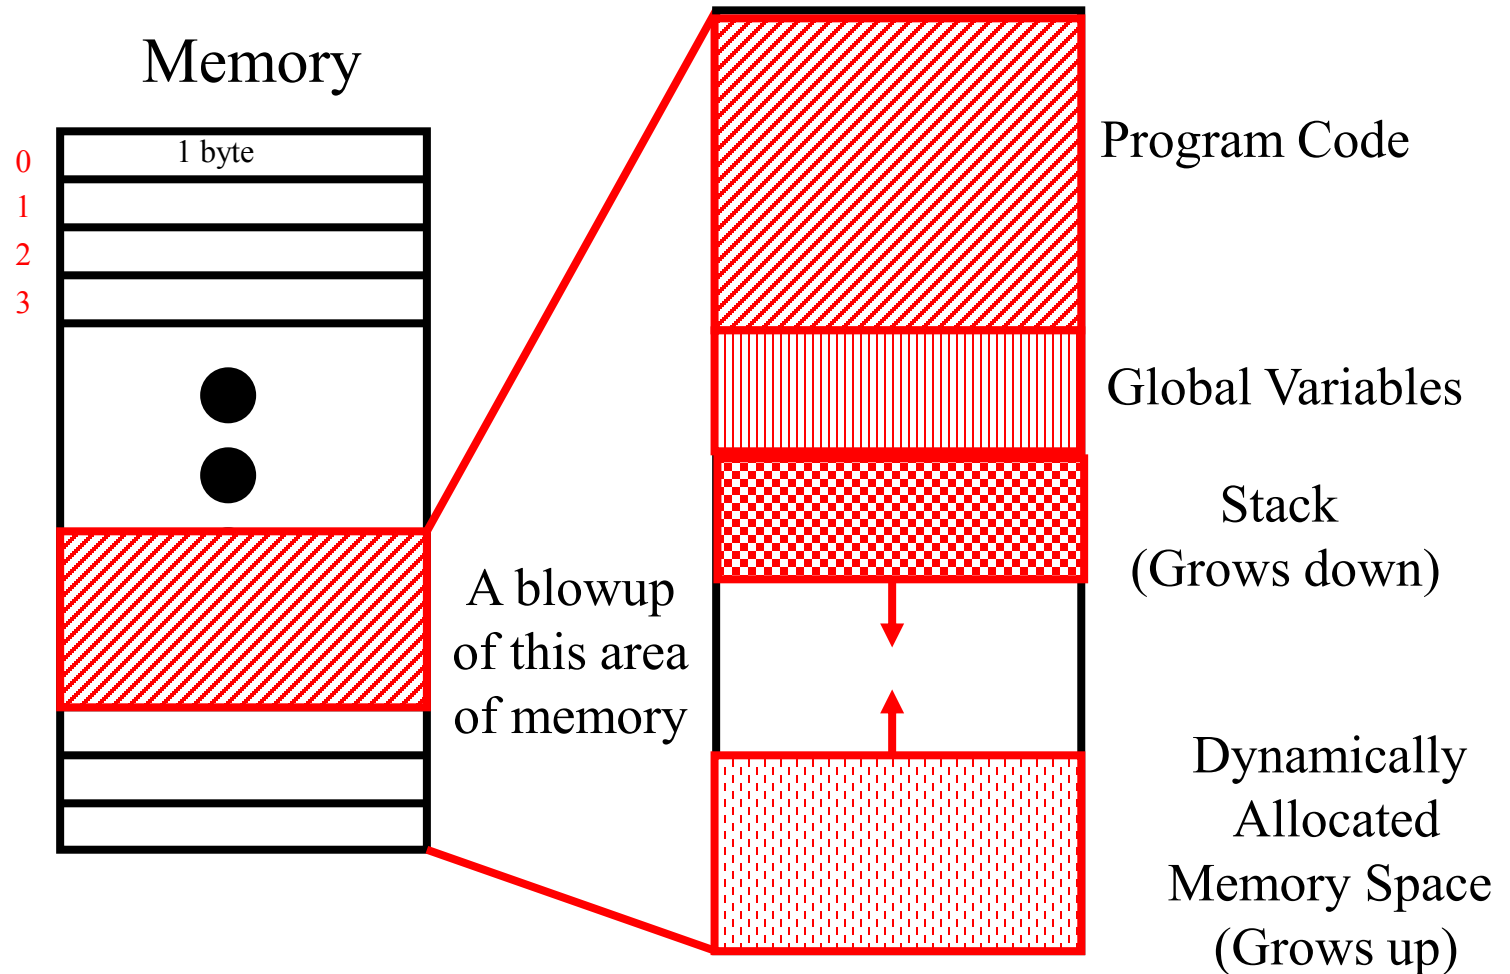


# Running a Program





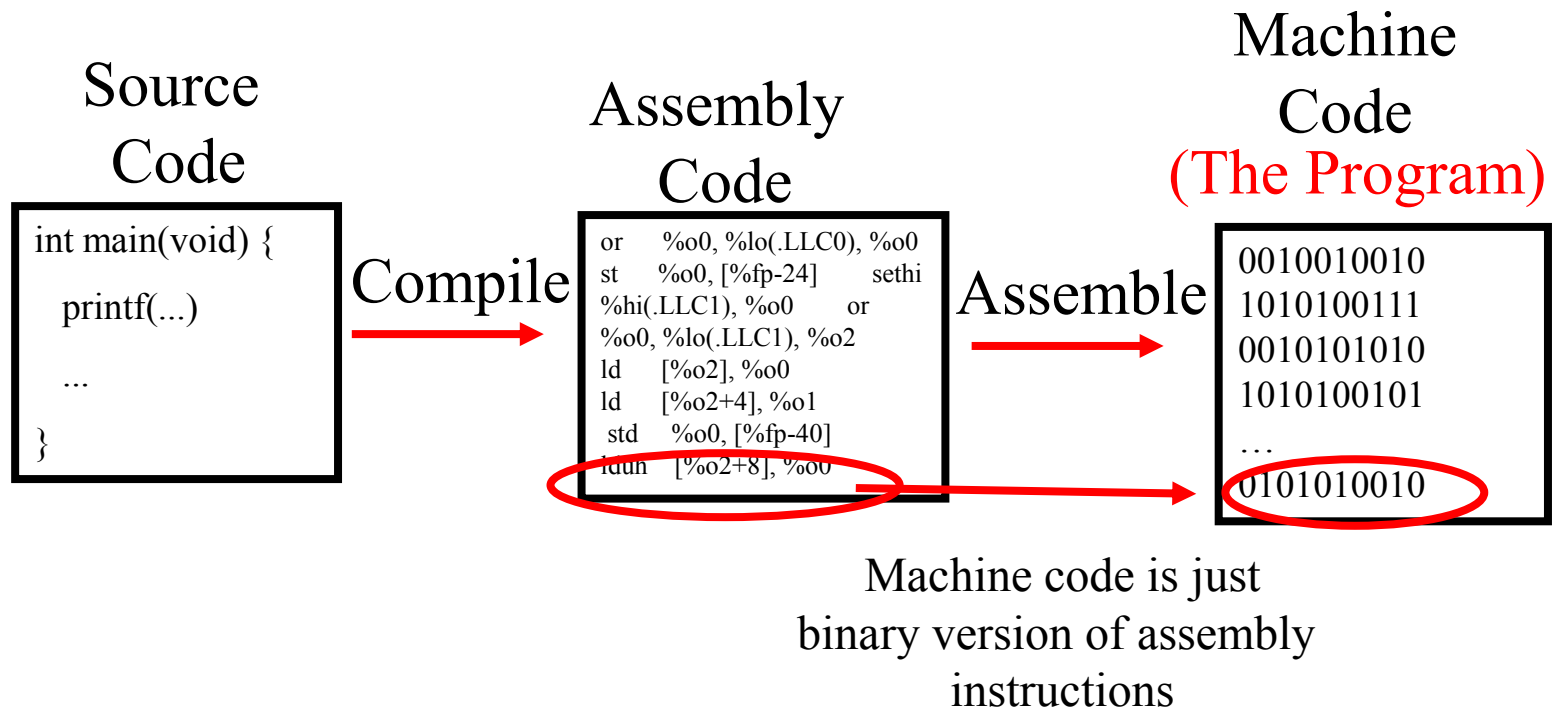
# Organization of Program in Memory



# Creating The Program Code

- How do we turn a high-level program (C++, Java) into something that the computer can run?

# Creating The Program Code – High Level



- By default, most compilers will both compile and assemble your source code.
- `gcc -S yourfile.c` will put the assembly code your program becomes into `yourfile.s`

# The compile/link process

- The C pre-processor expands macros
  - `#include`, `#define`, `#ifdef`, etc.
- The compiler parses your source, checks for errors and generates assembly code
- The compiler automatically calls the assembler, which converts assembly code to machine code
- If you are compiling an executable, the linker step tries to match function calls to function code (they might be in different files!)

# GCC

- The options you need to understand
  - `-g` Compile with debugging info for GDB
  - `-c` Compiles only, without linking (more later)
  - `-O3` Optimizes as much as possible
  - `-o` specifies the *name* of the output file
  - These should work with any Unix C/C++ compiler
- `-Wall`: a useful option for gcc
  - Turns on all warnings - which may point out potentially flawed code

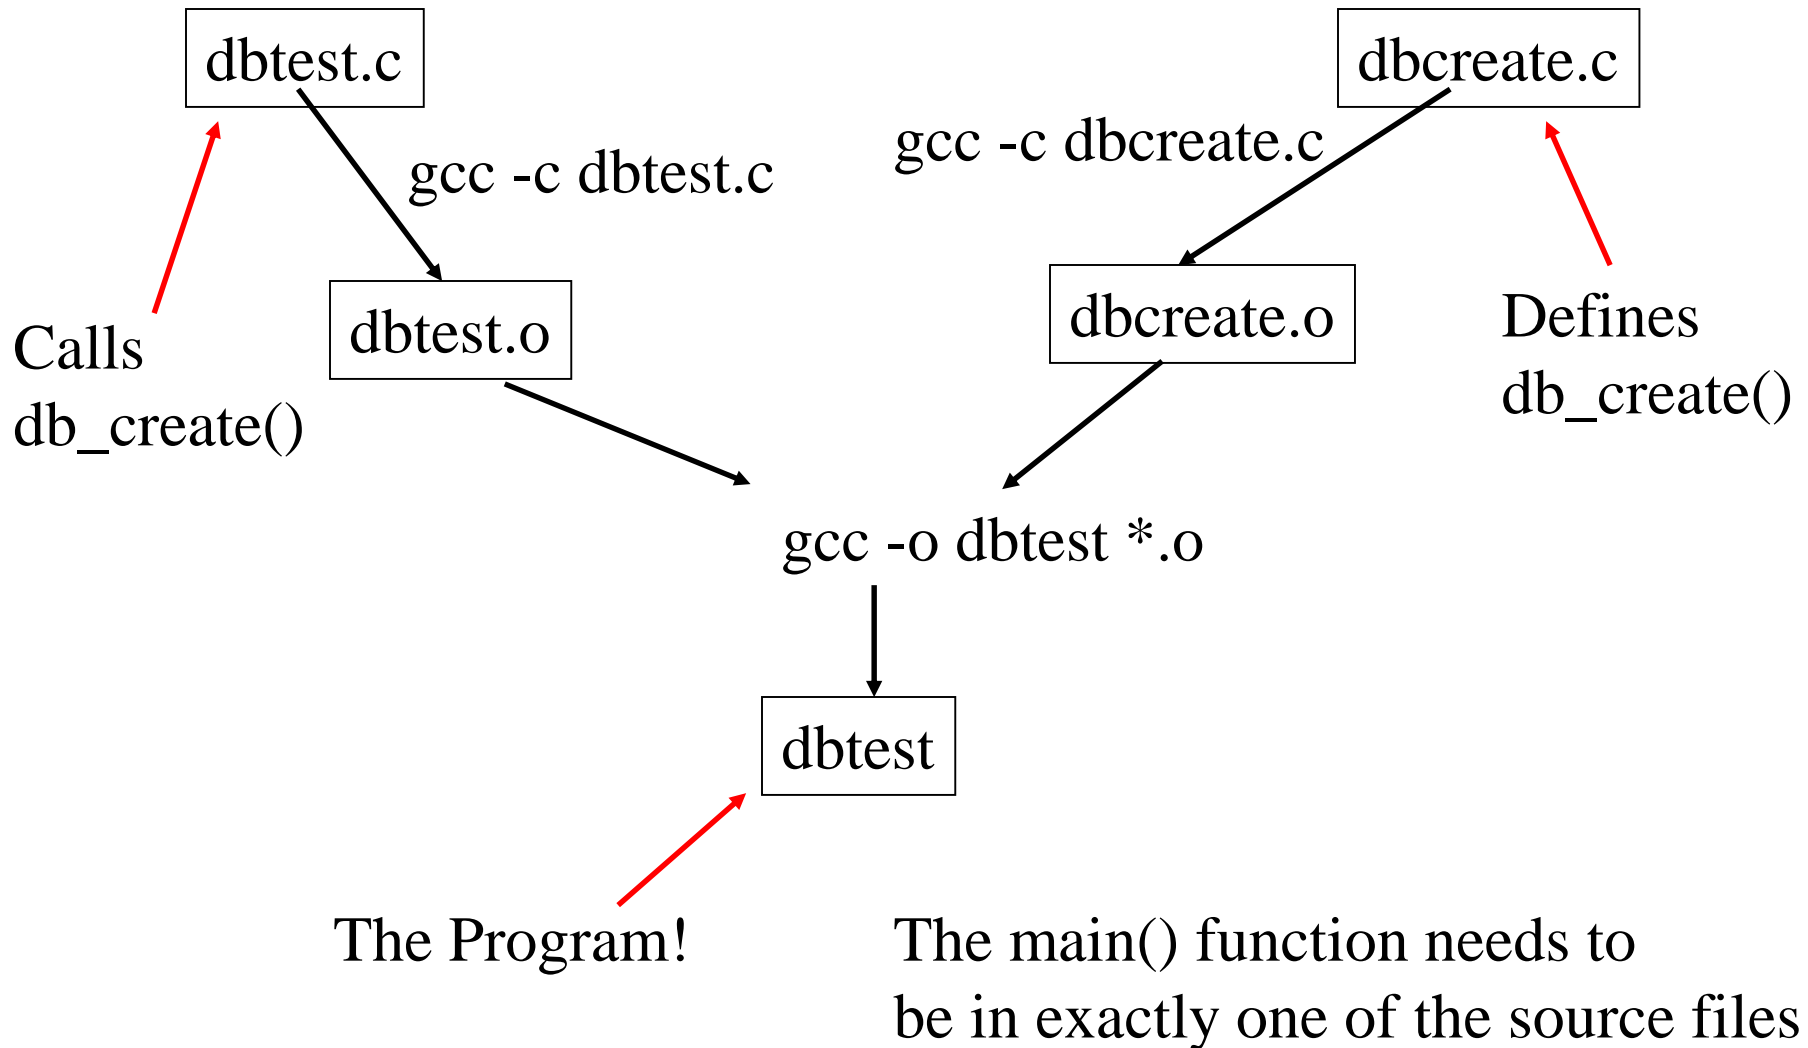
# Compiling an executable

- If you have one source (.c) file
  - `gcc -o dbtest dbtest.c`
    - Don't call your executable program `test...` why?
- If you have multiple source (.c) files
  - Option A: compile them all at once
    - `gcc -o dbtest dbtest.c dbcreate.c dbopen.c`
  - Option B: (more efficient) compile them one at a time without linking, then link them all together at the end

# Separate Compile and Link

- First compile all source files separately into object files (.o)
  - `gcc -c dbtest.c`
  - `gcc -c dbcreate.c`
  - `gcc -c dbopen.c`
  - `gcc -c dbread.c`
  - ...
- Once all the object (.o) files have been created, link them together to create an executable
  - `gcc -o dbtest dbtest.o dbcreate.o dbopen.o`

# Compile/Link





# Library Archives

- Library archives are collections of object files (.o) gathered into a single large file, with indexes to make accessing them fast
  - usually faster than having to read every .o file
  - easier to link with if you aren't changing the library object files frequently
- To create a library
  - First create all the object files (see previous slide)
  - Then use the `ar` command:
    - `ar -r libdb.a dbcreate.o dbopen.o dbread.o`

# Using Library Archives

- Include the library anywhere you can use an object file:  
`- gcc -o dbtest dbtest.o libdb.a`

# Concurrency

- Concurrency - multiple processes executing at the same time
- On UNIX concurrency is easy
  - Multiple processes can be running simultaneously
  - Multiple copies of the same program can be running
- Concurrency is very powerful
  - Greatly increases the efficiency of an OS:
    - While one process is waiting for I/O, another process can use the CPU

# Multi-user, multi-process system?

- Multiprogramming
  - More than one process can be ready to execute...
  - System calls trigger “context switches”: let the next process run
  - The process will not execute again until its system call returns
- Timesharing
  - CPU time split between multiple processes
    - Gives illusion that many processes are running at once



Can processes communicate? Yes! Stay tuned...

# Multi-user, multi-process system?

- What about multi-processor systems?
  - Each process can do multiprogramming AND timesharing

# Possible Complications

- Concurrently running processes can share data and/or resources
- What if multiple processes access the same resource at the same time?
- This is most likely a disaster
  - aka, “Race Condition”, “Oops”, or “aw carp”

# The Classic Example

- Two ATM machines
  - Each withdrawing \$20 from same account
- To update bank account balance
  - Read current balance into memory
  - Subtract \$20
  - Write new balance to the bank account

## Bank Balance


Process 1 reads the balance




Process 2 reads the balance



Process 1 subtracts and writes




Process 2 subtracts and writes





# Race conditions?

- Why are race conditions hard to detect?
  - They may only ever show up once
- Another way of saying it: “A race hazard (or race condition) is a flaw in a system or process where the output exhibits unexpected critical **dependence** on the relative timing of events.”



# Aw carp

- Most infamous race condition:
  - [http://en.wikipedia.org/wiki/Therac\\_25](http://en.wikipedia.org/wiki/Therac_25)
- People could outrun the system
  - the system was counting on a normally slow human, and didn't take into account people learning to use the system faster

# Lesson

- Concurrent update situation
  - 2+ processes accessing resource concurrently
  - At least one process might write
- **Must provide access control**
  - If one process is writing, no other process should access (read OR write) the resource
- "Locks" solve these problems
  - Only process owning lock may access (r/w) resource
  - Many ways to do locking (taught in other courses)
  - Locking usually requires support from the OS
    - But you can do it in software, too

# Lesson – use a lock file

```
do
{
    lock_fd = open(lock_file_path, O_WRONLY | O_CREAT | O_EXCL, 0644);

    if (lock_fd == -1)
    {
        if (errno == EEXIST)
        {
            // File already exists - wait a while, then try again
            sleep(1);
        }
        else
        {
            // An unexpected error - bail out
            perror("Couldn't open lock file\n");
            return(-1);
        }
    }
}
while (lock_fd == -1);
```

End