

Unix == Files

Benjamin Brewster

Adapted from slides by Jon Herlocker, OSU

Unix Paradigm

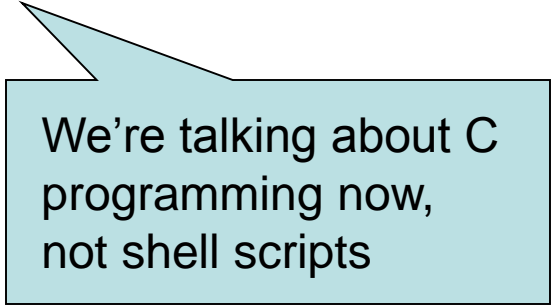
- Everything is a file
 - Even directories
 - Even hardware devices
 - Everything... except for processes

What is a File in UNIX?

- System Programmer View
 - Stream of bytes
 - An array
 - Newlines/carriage returns & tabs are all just bytes
 - Persistent
 - Usually stored on magnetic disk

Accessing a File in UNIX

```
# Open the file  
file_descriptor = open(...)
```



We're talking about C programming now, not shell scripts

```
# Call read(...) or write(...) any  
# number of times
```

```
# Close the file  
close(file_descriptor)
```

Opening a File

- File can be open for
 - read only - `O_RDONLY`
 - write only - `O_WRONLY`
 - read and write - `O_RDWR`
- When you open a file for writing...
 - Should you delete an existing file with the same name?
 - If not, where do you want to start writing
 - Beginning? End? Somewhere else?
 - If the file doesn't exist, should you create it?
 - If you create it, what should the initial access permissions be?

Open for Read

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <stdio.h>
int main(void)
{
    char file[] = "cs344/grades.txt";
    int file_descriptor;
    file_descriptor = open(file, O_RDONLY);
    if (file_descriptor < 0)
    {
        fprintf(stderr, "Could not open %s\n", file);
        exit(1);
    }
    exit(0);
}
```

Man pages again?

- In some cases there are two different man pages with the same name
 - Example :
 - “man read” will bring up the Bash manual page
 - “man -S 2 read” will bring up the read() system call page
- Make sure that you are looking at the right man page
 - Will say “Linux” at the bottom or top of *system call* pages (used in C). If it says Bash, it’s for the *shell* itself!
 - Man pages for *user programs* may say “GNU” or “FSF”
 - They **shouldn’t** say SunOS, Solaris, or Tcl!

Can't find the right man page?

- Use “man -k”, which does a keyword search of all man page titles
 - Example: “man -k pthread” returns all man pages with “pthread” in their title
 - “man -k split” shows you man pages related to “split”
- The results will show you which section of the manual a man page is found:

split	(1)	- split a file into pieces
split	(n)	- Split a string into a proper Tcl list
splitdiff	(1)	- separate out incremental patches
tiffsplit	(1)	- split a multi-image TIFF into single-image TIFF files
...		
- You can then specify the right man page section
 - Example: “man -S n split”

Open for Write

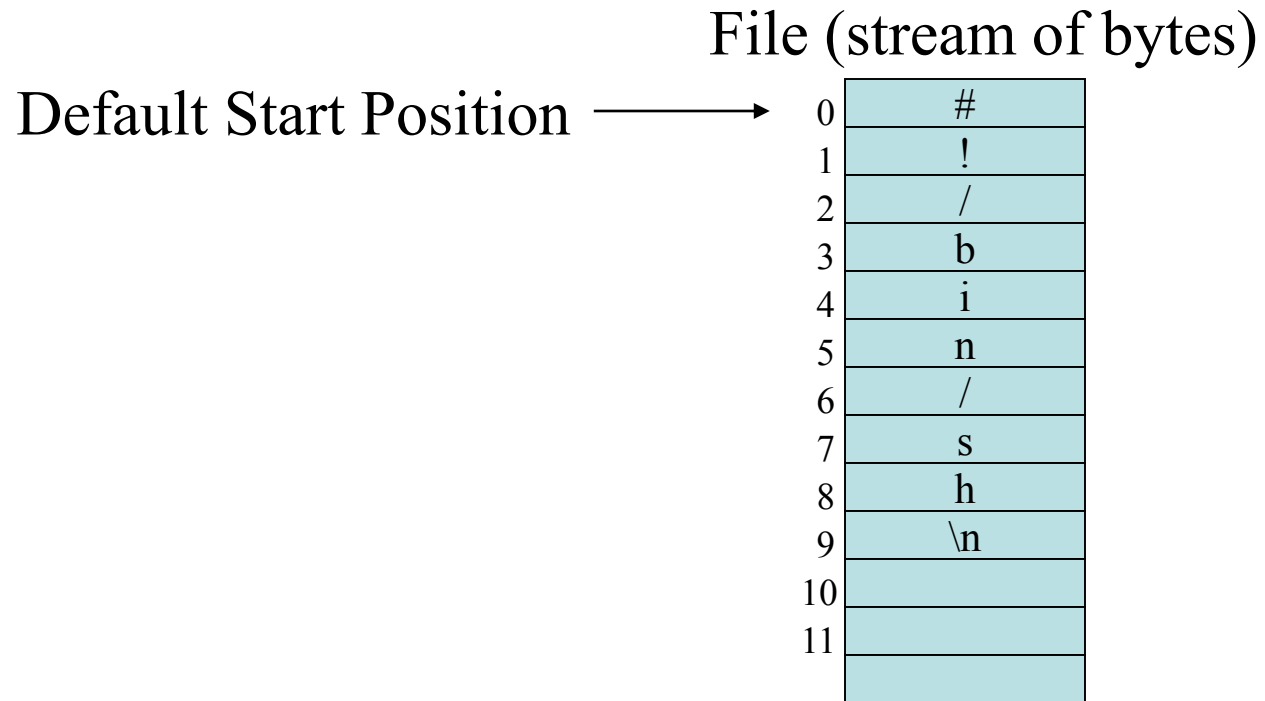
```
int main(void)
{
    char file[] = "cs344/grades.txt";
    int file_descriptor;
    file_descriptor = open(file, O_WRONLY);
    if (file_descriptor < 0)
    {
        fprintf(stderr, "Could not open %s\n", file);
        perror("in main");
        exit(1);
    }
    exit(0);
}
```

```
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
int main(void)
{
    int file_descriptor;
    char *file = "test.txt";
    ssize_t nread;
    ssize_t nwritten;
    char buffer[512];
    file_descriptor = open(file, O_RDWR);
    if (file_descriptor == -1)
        exit(1);
    nread = read(file_descriptor, buffer, 512);
    nwritten = write(file_descriptor, buffer, nread);
    exit(0);
}
```



Read/Write example

The file pointer



Truncating an Existing File

- When you open a file for writing
 - Should you delete an existing file with the same name?
 - To delete it: `O_TRUNC`
 - Example:
 - `file_descriptor = open(file, O_WRONLY | O_TRUNC);`
 - Deletes all data in existing file
 - Sets write pointer to position 0

Appending to an Existing File

- O_APPEND

- Example

- `file_descriptor = open(file, O_WRONLY | O_APPEND);`

- Before every write, the file pointer will be reset to the end of the file

The file pointer

File (stream of bytes)

0	#
1	!
2	/
3	b
4	i
5	n
6	/
7	s
8	h
9	\n
10	
11	
⋮	

O_APPEND →

Creating a New File

- O_CREAT
- Flag indicates that the file should be created if it doesn't already exist
- Example
 - `file_descriptor = open(file, O_WRONLY | O_CREAT);`

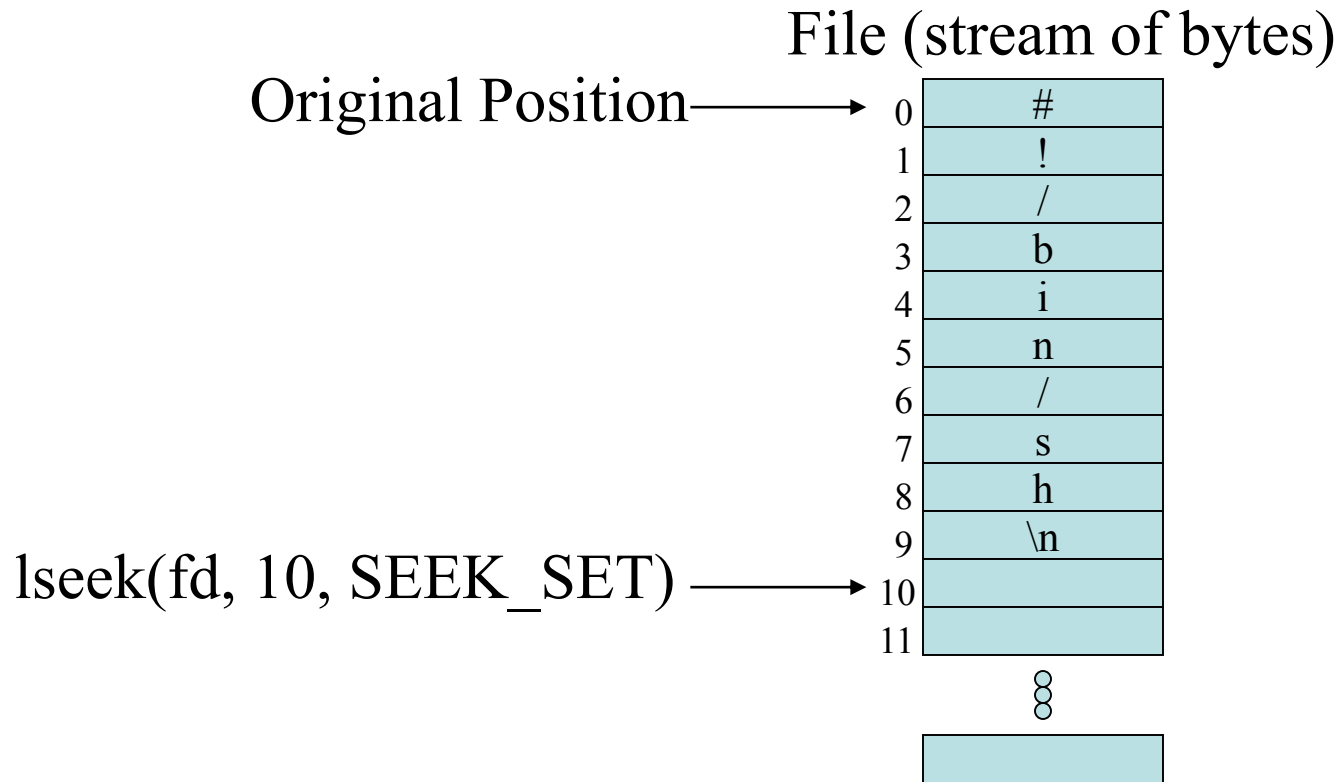
Access Permissions

- Third parameter only applies when new file is created (i.e. using O_CREAT)
- Third parameter is octal number
 - bits of the octal number signify the permissions
- If you don't specify it, there are reasonable defaults
- Example:
 - `file_descriptor = open(file, O_WRONLY|O_CREAT, 0644);`
- See "man chmod"

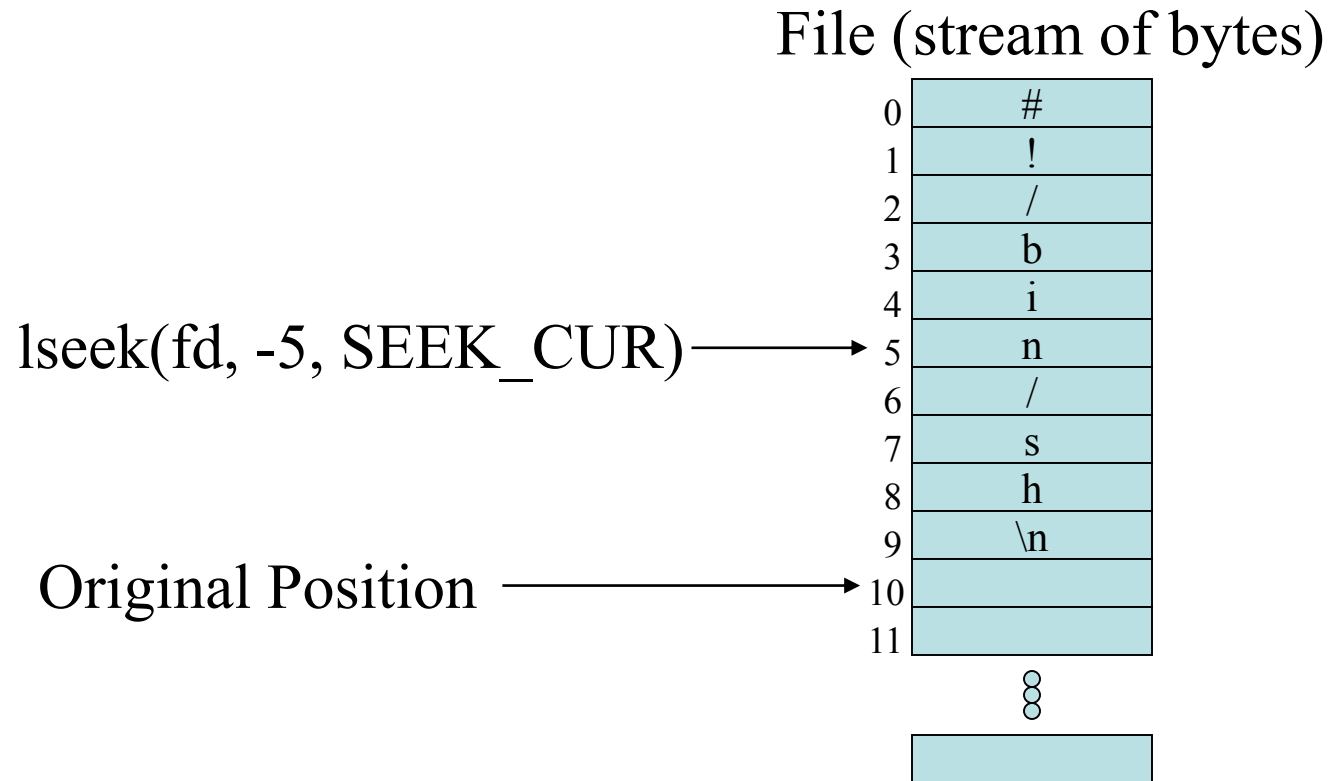
lseek

- Random access
 - Jump to any byte in a file
- Move to byte #16
 - `newpos = lseek(file_descriptor, 16, SEEK_SET);`
- Move forward 4 bytes
 - `newpos = lseek(file_descriptor, 4, SEEK_CUR);`
- Move to 8 bytes from the end
 - `newpos = lseek(file_descriptor, -8, SEEK_END);`

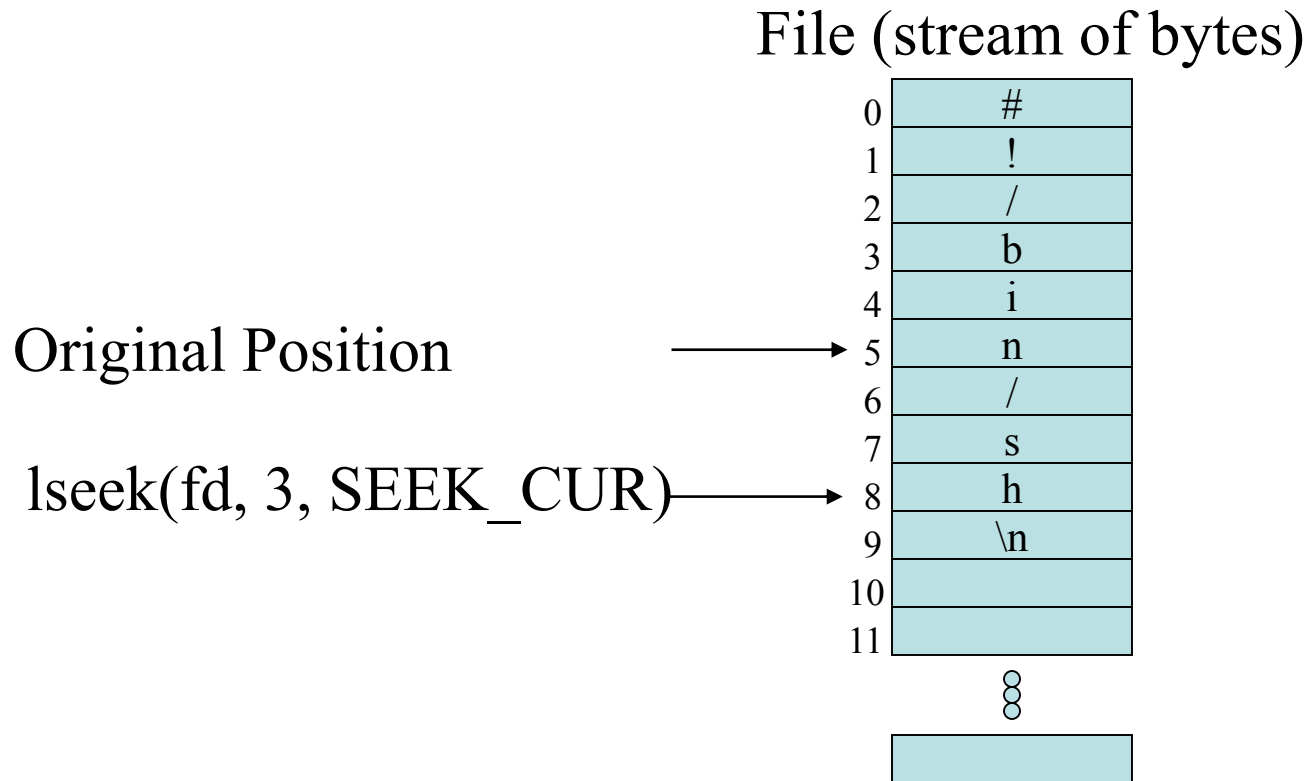
lseek - SEEK_SET



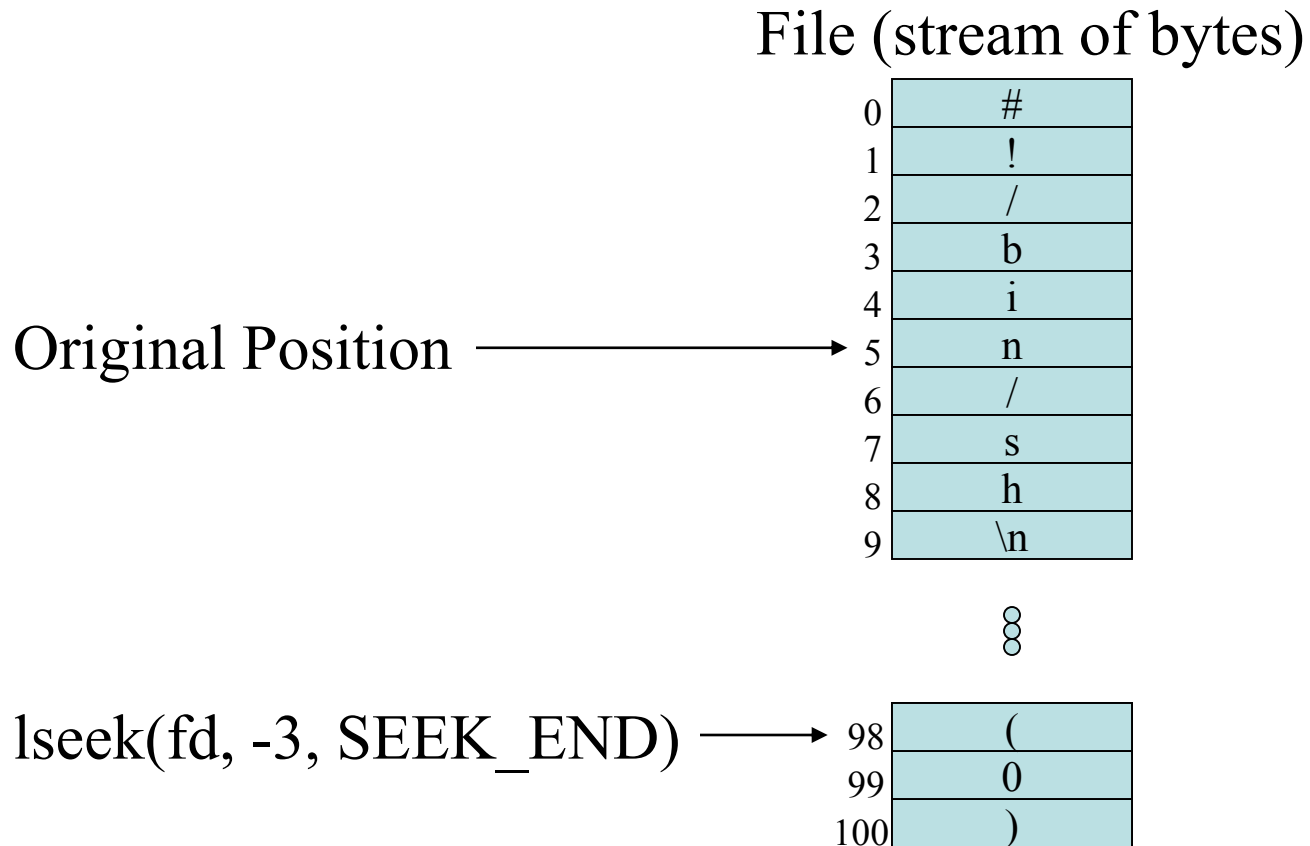
lseek - SEEK_CUR



lseek - SEEK_CUR




lseek - SEEK_CUR



Read/Write and the File Pointer

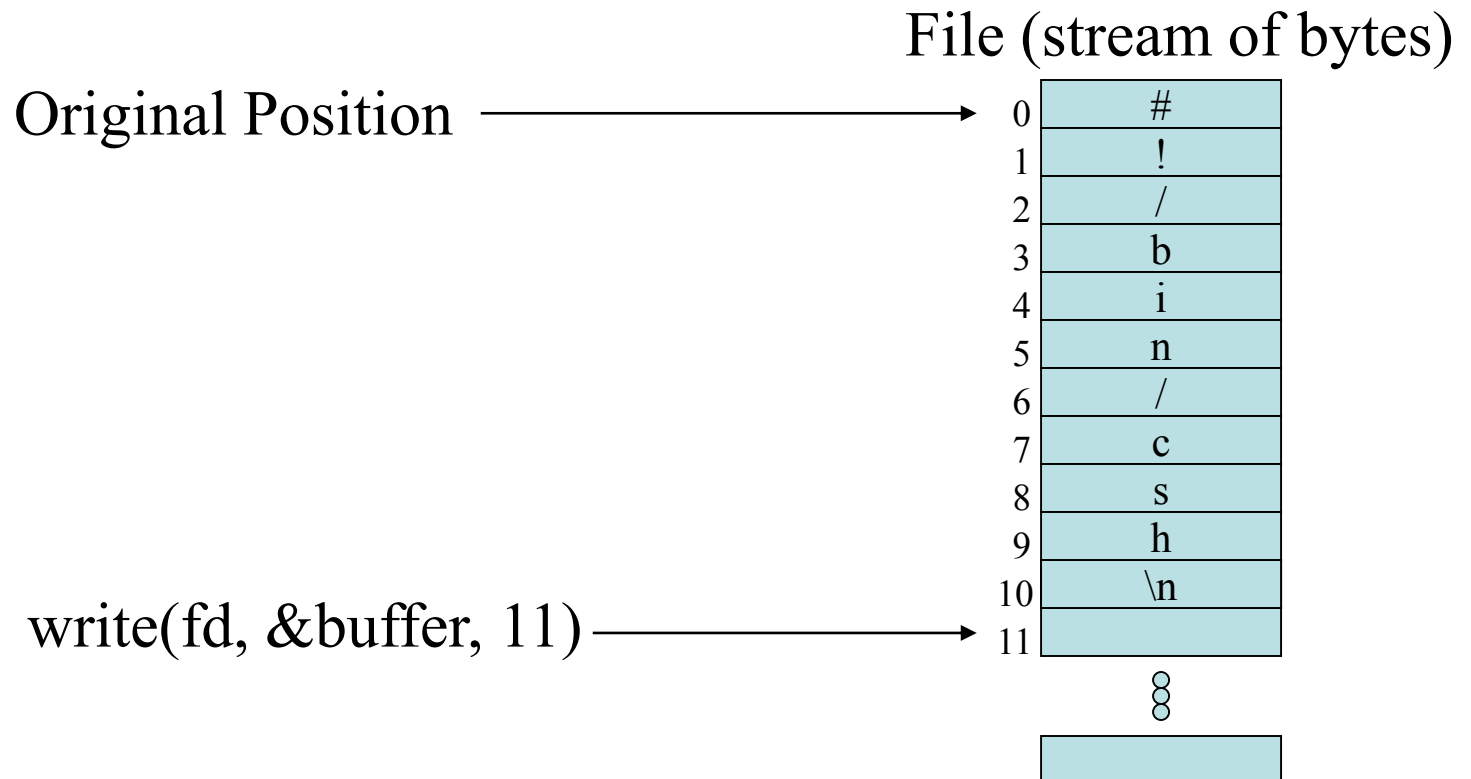
- Both `read()` and `write()` will change the file pointer.
- The pointer will be incremented by exactly the number of bytes read or written.

read



write

buffer: `#!/bin/csh`



The Standard IO Library

- fopen, fclose, printf, fprintf, sprintf, scanf, fscanf,getc, putc, gets, fgets, etc.
- #include <stdio.h>

Why use stdio library

- Automatically buffers input and output intelligently
- Easy to work in line mode
 - ie read one line at a time
 - write one line at a time
- Powerful string and number formatting

Why use read()/write()



As opposed to stdio

- Maximal performance
 - IF you know exactly what you are doing
 - No additional hidden overhead from stdio
- Control exactly what is written/read at what times

Some stdio functions

- `fclose` Close a stream.
- `feof` Check if End Of File has been reached.
- `fgetc` Get next character from a stream.
- `fgetpos` Get position in a stream.
- `fopen` Open a file.
- `fprintf` Print formatted data to a stream.
- `fputc` Write character to a stream.
- `fread` Read block of data from a stream.
- `fseek` Reposition stream's position indicator.
- `getc` Get the next character.
- `getchar` Get the next character from stdin.

Some more stdio functions

- `gets` Get a string from stdin.
- `printf` Print formatted data to stdout.
- `putc` Write character to a stream.
- `putw` Write an integer to a stream.
- `remove` Delete a file.
- `rename` Rename a file or directory.
- `rewind` Reposition file pointer to the beg. of a stream.
- `scanf` Read formatted data from stdin.
- `sprintf` Format data to a string.
- `sscanf` Read formatted data from a string.
- `ungetc` Push a character back into stream.

Obtaining File Information

- `stat()` and `fstat()`
- Retrieve all sorts of information about a file
 - Which device it is stored on
 - Ownership/Permissions of that file
 - Number of links
 - Size of the file
 - Date/Time of last modification and access
 - Ideal block size for I/O to this file

```
struct stat statbuf;
int r;

r = stat(argv[1], &statbuf);
if (r == -1)
{
    fprintf(stderr, "Could not stat %s\n", argv[1]);
    perror(argv[0]);
    exit(1);
}

/* print out the size of the file */
printf("The logical size of %s is %ld bytes\n",
    argv[1], statbuf.st_size);
printf("The physical size of %s is %ld bytes\n",
    argv[1], statbuf.st_blocks * 512);
```

```
struct stat statbuf;
int r;
int fd;

fd = open(argv[1], O_RDONLY);
if (fd == -1) { /* handle error */  exit(1); }

r = fstat(fd, &statbuf);
if (r == -1)
{
    fprintf(stderr, "Could not fstat %s\n", argv[1]);
    perror(argv[0]);
    exit(1);
}

printf("The logical size of %s is %ld bytes\n", argv[1],
    statbuf.st_size);
printf("The phys      size of %s is %ld bytes\n", argv[1],
    statbuf.st_blocks * 512);
```


End