

Pipes and Redirection

It just never stops

Benjamin Brewster

Adapted from slides by Jon Herlocker, OSU

Sharing open files

- When you exec, you replace your current process with a new one
- But your files are still open
 - This may not be what you want
- Are there other ways to share files?

Sharing open files

- Preventing open files from being shared across an exec
 - close-on-exec
- I/O redirection
 - Redirecting input to and from files on disk
 - Pipes: redirecting input and output between different processes
 - akin to command line piping, but not the same

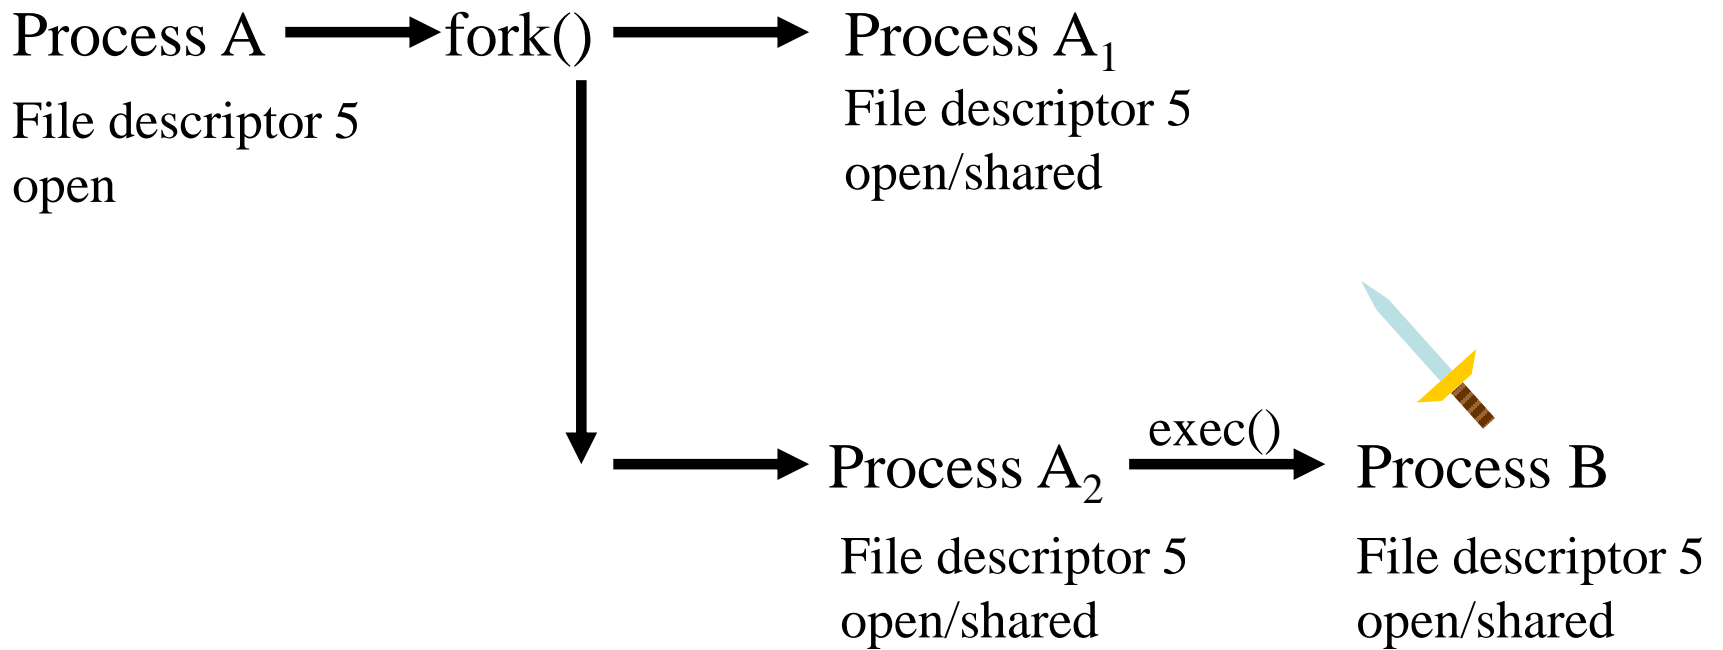
Close on exec

- Tells the kernel to close open files on exec
 - why would we want to do this?
- Open files are "inherited" by child processes
 - Thus the file pointer is shared (!)
 - Security/sensitive data

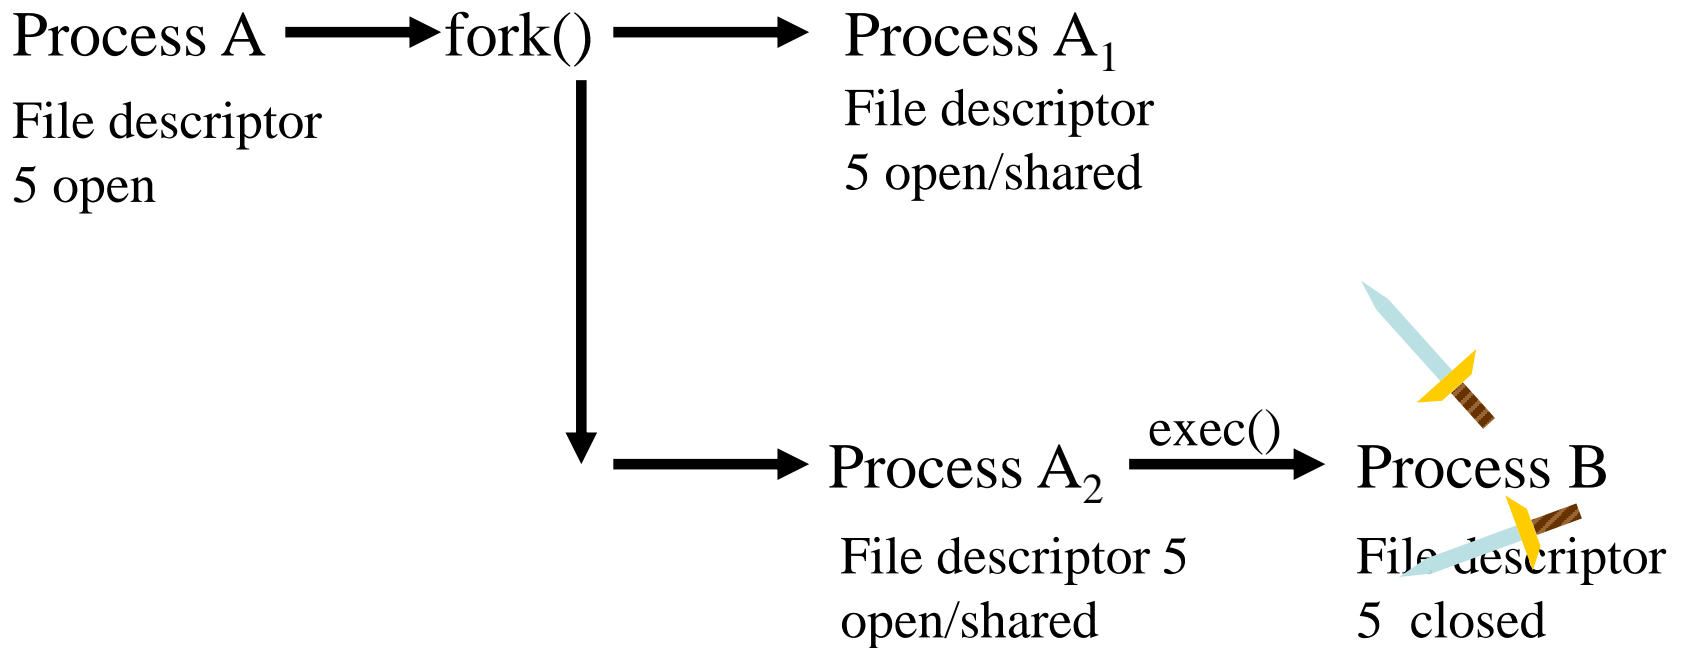
Close on exec

- You set the "close-on-exec" flag for every file descriptor you do not want to share
- The "close-on-exec" flag is inherited through fork
 - So you can set the close-on-exec flag in the parent and if the child does an exec() the file will be closed, as well

Normally



With close on exec



Close on exec example

```
#include <fcntl.h>
```

```
...
```

```
int fd;
```

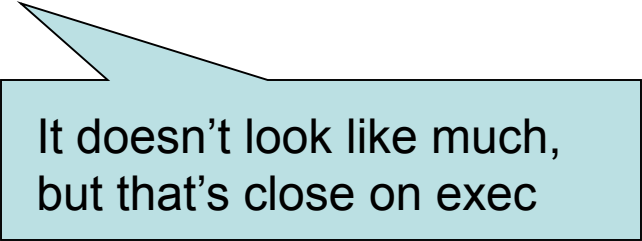
```
fd = open("file", O_RDONLY);
```

```
...
```

```
fcntl(fd, FD_SETFD, 1);
```

```
...
```

```
exec...
```



It doesn't look like much,
but that's close on exec

I/O redirection

- We saw I/O redirection in the shell
 - `ls > file`
 - `stats < file1`
 - `cat longfile | more`
 - `find . -name paper -print 2> /dev/null`
 - `echo "an error occurred" 1>&2`
- I/O redirection is possible *because* open files are shared across `fork()` and `exec()`

Important background

- The kernel opens stdin, stdout, and stderr automatically for every process created
- File descriptor 0 is stdin
- File descriptor 1 is stdout
- File descriptor 2 is stderr
- They default to reading and writing to the terminal

I/O redirection

- The trick: you can change where the standard I/O streams are coming from and/or going *after* the fork *but before* the exec

Redirecting stdout

Process starts



opens new output file - fd = 3




Duplicates fd 3 into fd 1
(overwrites fd 1)




Calls exec

Redirecting I/O



Redirecting stdout


1. First open the new file



Redirecting stdout

2. Call `dup2 ()` to change fd 1 to point where fd 3 points:

```
dup2 (3, 1);
```



```
fd = open("myout.data", O_WRONLY|O_CREAT|O_TRUNC, 0644);
if (fd == -1)
{
    perror("open");
    exit(1);
}

fd2 = dup2(fd, 1);
if (fd2 == -1)
{
    perror("dup2");
    exit(2);
}

execlp(prog1, prog1, NULL);
perror("exec");
exit(3);
```

Change where stdout is pointing:

- Make it point to where our newly opened output file points


```
fd = open("myin.data", O_RDONLY);
if (fd == -1)
{
    perror("open");
    exit(1);
}

fd2 = dup2(fd, 0);
if (fd2 == -1)
{
    perror("dup2");
    exit(2);
}

execlp(prog1, prog1, NULL);
perror("exec");
exit(3);
```

Change where stdin is pointing:

- Make it point to where our newly opened input file points

Real Inter-Process Communication (IPC)


- IPC methods in UNIX
 - **Intermediate/temporary files** (often together with I/O redirection)
 - **Pipes** - IPC between two processes forked by a common ancestor process
 - **FIFOs** (named pipes) - communication between any two processes on the same machine
 - **SysV message IPC** - communication between any two processes on the same machine
 - Not a simple byte stream
 - Supports message categories - often used for priorities
 - **Sockets** - communication between any two processes potentially separated by a network

Between process IPC

- I/O redirection with `dup2 ()` allows you to redirect input and output between processes and files
- How do we redirect input between processes and other processes on the same machine?
 - Use temporary/intermediate files
 - Writes to disk are slow
 - No good way to track when the other process is ready to receive or send new data
 - Better answer: use **pipes**!

Pipes

- Pipes provide a way to connect an output-only file descriptor in one process to an input-only file descriptor in another process



Creating a Pipe

- Pipes are possible because file descriptors are shared across `fork()` and `exec()`
- A parent process creates a pipe
 - Results in two new open file descriptors, one for input and one for output
- The parent process forks (and possibly execs)
 - Parent and child have the fds created with the pipe
- The child process now reads from the input file descriptor, and the parent process writes to the output file descriptor
 - or vice-versa

The `pipe()` function

- You pass `pipe()` an array of two integers, where it stores the two new open file descriptors
- The first is the input file descriptor, and the second is the output file descriptor
- One of the descriptors is used by the parent process and the other is used by the child process

```
int r, pipeFDs[2];
char message[512];
pid_t spawnpid;

...

if (pipe(pipeFDs[2] == -1)
{
    perror("Hull Breach!");
    exit(1);
}

spawnpid = fork();
switch (spawnpid)
{
    case 0: // Child
        close(pipeFDs[0]); // close the input file descriptor
        write(pipeFDs[1], "hi parent, this is the child!!", 41);
        exit(0);

    default: // parent
        close(pipeFDs[1]); // close output file descriptor
        r = read(pipeFDs[0], message, sizeof(message));
        if (r > 0)
            printf("Message received from child: %s\n", message);
        exit(0);
}
```

flow control with `read()`

- `read()` succeeds if data is available
 - Receives the data and returns immediately
 - The return value of `read()` tells you how many bytes were read - *it may be less than you requested*
- if data is not available, `read()` will block waiting for data (your process execution is suspended until data arrives)
 - `read()` is a system call

flow control with `write()`

- Write will not return until all the data has been written
 - `write()` is a system call
- Pipes have a certain size
 - Only so much data will fit in a pipe
 - If the pipe fills up, and there is no more room, `write()` will block until space becomes available (ie somebody `reads` the data from the pipe)

Pipe Recap

- `write()` puts bytes in the pipe, `read()` takes them out
- It is possible to determine the size of the pipe - see `fpathconf()`



Programming note

- Checking the return value of `read()` becomes very important
 - Not just if return value is -1 (an error)
 - The return value will tell you if the desired number of bytes was not read
- Same goes for `write()`

Closing Pipes

- What happens if a process closes their end of the pipe when the other process is still trying to read or write to the pipe?



Closing Pipes

- Process A closes output pipe.
 - If process B is currently blocked on a `read()`, then process B's `read()` will return 0
- Process B closes input pipe
 - If process A tries to write to the pipe, `write()` will return -1, and `errno` (in process A) will be set to `EPIPE`
 - Process A will be sent the `SIGPIPE` signal

