


UNIX Networking 1

Benjamin Brewster

Adapted from Jon Herlocker, OSU

Network Communication : IPC




- Terminology
 - **Client/Server architecture** - a networking arrangement such that one process (server) is continuously waiting for new connections from other processes (clients)
 - **Client** - process that initiates the connection
 - **Server** - process that is always running, waiting for new connection initiations from client processes

UNIX Daemons – can be server

- A daemon is a process running in the background – a background process
 - `syslogd` maintains the system log
- In UNIX, daemons all have the `init` process (`pid == 1`) as their parent

Client/Server Example: Web Server



- The web server is always running and looking for new connections
- Potentially unknown number of clients who may connect at any time
 - Netscape, IE, Lynx, etc.

TCP/IP

- TCP/IP provides communication between two processes, potentially separated by a network
- TCP/IP stands for Transmission Control Protocol / Internet Protocol
- TCP is what your application interacts with, and IP provides the addressing system for sending

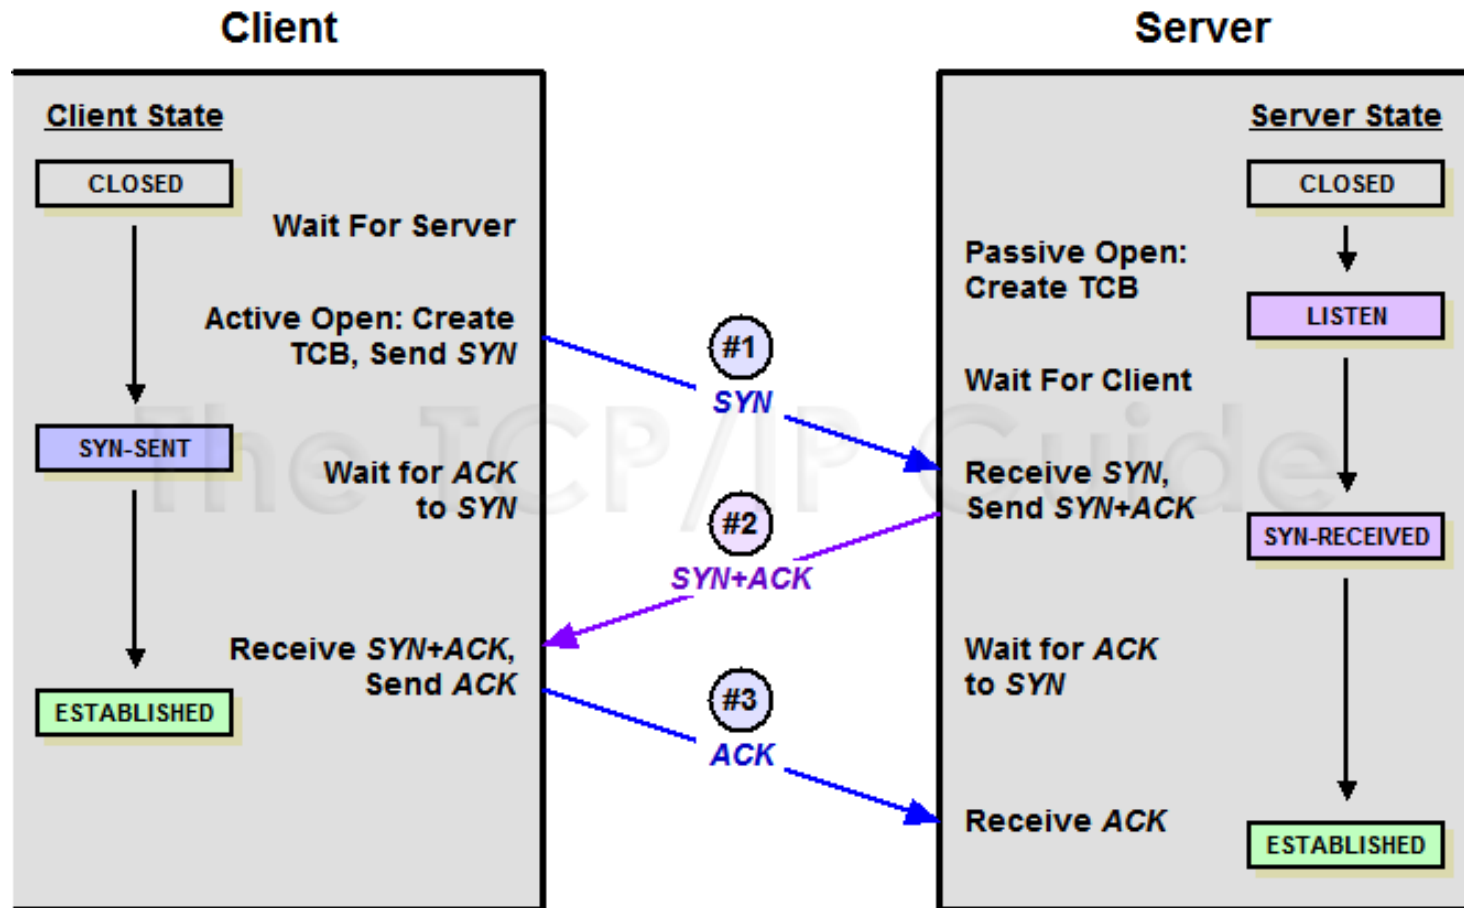
What is TCP/IP

- A transmission protocol
- Provided by the OS to user processes
- Only passes bytes between processes
 - It does not interpret those bytes
- You need an "application protocol" on top of TCP/IP which defines what the bytes mean
 - This application protocol is what your program does with the bytes it receives

Application Protocols

- TCP/IP is like a phone connection
 - A phone transfers sound between two people
 - A phone does not interpret the meaning of the sound
- The telephone application protocol (A calling B)
 - B says "hello" (First step)
 - A responds "is so-and-so available" (Second Step)
 - B responds "yes", "no", or "speaking" (Third Step)
 - If B is the desired person, then begin talking
 - If B is not available, then leave message and hang up

TCP Handshaking



Example: Web server application protocol

- **HTTP** - hyper text transfer protocol
- Standardized protocol defining:
 - how client requests are formatted
 - how server responses are formatted
- The protocol is text-based
 - All requests, responses, and errors
- Relatively simple protocol
 - Primarily used to support one feature: requesting a file to be downloaded

HTTP Protocol

- Simple request mode
 - Client connects to server then sends:
 - GET abc.html
 - Server receives and parses the get command and returns the file requested
- Enhanced request mode
 - Client connects to the server then sends:
 - GET index.html HTTP/1.0
 - Server responds with some header information, such as server type and version, then a blank line, then the data file

Text Protocol Debugging Tool

- `telnet` helps debug text-based protocols
- We're not using telnet here for shell access – OSU requires that we use only the SSH protocol
- You can pass telnet a second parameter that specifies the port you want to connect to
 - To connect to a web server, we typically use port 80
 - `telnet eecs.oregonstate.edu 80`

Demo of telnet and HTTP

1. `% telnet eecs.oregonstate.edu 80`
2. `<web server> GET / HTTP/1.1`
3. `<web server> Host: eecs.oregonstate.edu`
4. `<web server> (Enter twice)`

Non-text-based Application Protocols

- Not all application protocols are text-based
- TCP/IP has no problem transferring binary data
- Advantages of text-based protocols
 - Easy to debug
 - Easy to communicate and understand (and teach!)
- Disadvantages of text-based protocols
 - Not very compact or efficient
 - Server can spend a lot of time just parsing text
 - Very important for text protocols to be simple!

Internet network layers

Application Protocol (ie HTTP)
TCP or UDP
IP
Ethernet

- **TCP** - Transmission Control Protocol - connection-oriented, guaranteed data transport
- **UDP** - Universal Datagram Protocol - connectionless, not guaranteed
- **IP** - Internet Protocol - naming, addressing and routing, independent of physical connection type
- **Ethernet** - the addressing and signaling protocol used across a copper cable connecting a collection of machines

TCP versus IP

- Once the TCP layer sends out packets, its up to IP to get them where they are going
- Problem: IP does not guarantee
 - Data integrity
 - Packet order
 - Prevention of duplicates
 - Packet will actually arrive

TCP versus IP

- TCP, however, can detect if IP is having trouble, and can
 - Re-order packets
 - Request packet re-transmission
 - Drop duplicate packets

TCP

- Most commonly used protocol for transferring information
- Provides a byte stream interface (like stdio)
 - Data is guaranteed to not be lost, even over networks where packets can be dropped
 - Data arrives in the order it was sent
 - Connection oriented - each side of the connection maintains resources to keep the connection open until it is explicitly closed
 - A TCP connection is bi-directional - traffic can be sent across the connection in either direction

UDP

- Used less frequently than TCP
- Provides a very different interface
 - Connectionless (No handshaking, etc.)
 - Data is broken into packets called datagrams
 - Server does not remember clients between datagrams
 - Datagrams may be dropped by the network
 - Datagrams may arrive out of order

TCP Benefits

- Benefits of TCP (over UDP)
 - Error-free data transfer
 - Ordered-data transfer
 - Retransmission of lost packets
 - Discarding duplicate packets
 - Congestion throttling

UDP Benefits

- UDP has much less overhead than TCP
- When to use UDP
 - Streaming video/audio
 - Mass broadcasting
 - Asynchronous communication
 - GAMEZ

Internet Protocol (IP)

- IP specifies:
 - How we address machines on the network
 - If the machine we are addressing is not on the same local network: how the data is routed
- Each network interface (network card) has an **IP address**, which must be unique
- IP(v4) address are 32 bit numbers, but are usually *represented* as four one-byte numbers, separated by periods:
 - Ex: 128.101.34.200
 - Ex: 128.193.55.70

IP 4 vs. 6

- IP version 4 (IPv4) uses 32-bit addresses, and therefore has $2^{32} = 4,294,967,296 = 4 \times 10^9$ unique addresses
 - This is not enough for the world - in my house alone:
 - Xbox One, Xbox 360, my laptop, wife's laptop, my phone, wife's phone, printer
 - A lot of this can be handled by routers, which create your own LAN, but this shows the problem

IP 4 vs. 6

- IPv6 uses 128-bit addresses
 - $2^{128} > 2^{32}$
 - $3.4 * 10^{38} > 4.2 * 10^9$
 - This is **50 octillion** addresses for **each** of 6.5 billion people on earth

50 octillion addys*:

50,000,000,000,000,000,000,000,000,000,000,000,000,000